

Importing Libraries

```
In [109]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

```
In [110]: df=pd.read_csv(r"C:\Users\user\Desktop\csvs_per_year\csvs_per_year\madrid_2004.
df
```

Out[110]:

| | date | BEN | CO | EBE | MXV | NMHC | NO_2 | NOx | OXY | O_3 | PI |
|--------|---------------------|------|------|------|------|------|------------|------------|------|-----------|--------|
| 0 | 2004-08-01 01:00:00 | NaN | 0.66 | NaN | NaN | NaN | 89.550003 | 118.900002 | NaN | 40.020000 | 39.990 |
| 1 | 2004-08-01 01:00:00 | 2.66 | 0.54 | 2.99 | 6.08 | 0.18 | 51.799999 | 53.860001 | 3.28 | 51.689999 | 22.950 |
| 2 | 2004-08-01 01:00:00 | NaN | 1.02 | NaN | NaN | NaN | 93.389999 | 138.600006 | NaN | 20.860001 | 49.480 |
| 3 | 2004-08-01 01:00:00 | NaN | 0.53 | NaN | NaN | NaN | 87.290001 | 105.000000 | NaN | 36.730000 | 31.070 |
| 4 | 2004-08-01 01:00:00 | NaN | 0.17 | NaN | NaN | NaN | 34.910000 | 35.349998 | NaN | 86.269997 | 54.080 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 245491 | 2004-06-01 00:00:00 | 0.75 | 0.21 | 0.85 | 1.55 | 0.07 | 59.580002 | 64.389999 | 0.66 | 33.029999 | 30.900 |
| 245492 | 2004-06-01 00:00:00 | 2.49 | 0.75 | 2.44 | 4.57 | NaN | 97.139999 | 146.899994 | 2.34 | 7.740000 | 37.689 |
| 245493 | 2004-06-01 00:00:00 | NaN | NaN | NaN | NaN | 0.13 | 102.699997 | 132.600006 | NaN | 17.809999 | 22.840 |
| 245494 | 2004-06-01 00:00:00 | NaN | NaN | NaN | NaN | 0.09 | 82.599998 | 102.599998 | NaN | NaN | 45.630 |
| 245495 | 2004-06-01 00:00:00 | 3.01 | 0.67 | 2.78 | 5.12 | 0.20 | 92.550003 | 141.000000 | 2.60 | 11.460000 | 24.389 |

245496 rows × 17 columns

Data Cleaning and Data Preprocessing

```
In [111]: df=df.dropna()
```

```
In [112]: df.columns
```

```
Out[112]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
                'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
               dtype='object')
```

```
In [113]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 19397 entries, 5 to 245495  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   date        19397 non-null  object  
1   BEN         19397 non-null  float64  
2   CO          19397 non-null  float64  
3   EBE         19397 non-null  float64  
4   MXY         19397 non-null  float64  
5   NMHC        19397 non-null  float64  
6   NO_2        19397 non-null  float64  
7   NOx         19397 non-null  float64  
8   OXY         19397 non-null  float64  
9   O_3         19397 non-null  float64  
10  PM10        19397 non-null  float64  
11  PM25        19397 non-null  float64  
12  PXY         19397 non-null  float64  
13  SO_2        19397 non-null  float64  
14  TCH         19397 non-null  float64  
15  TOL         19397 non-null  float64  
16  station     19397 non-null  int64  
dtypes: float64(15), int64(1), object(1)  
memory usage: 2.7+ MB
```

```
In [114]: data=df[['CO' , 'station']]
data
```

Out[114]:


| | CO | station |
|--------|------|----------|
| 5 | 0.63 | 28079006 |
| 22 | 0.36 | 28079024 |
| 26 | 0.46 | 28079099 |
| 32 | 0.67 | 28079006 |
| 49 | 0.30 | 28079024 |
| ... | ... | ... |
| 245463 | 0.08 | 28079024 |
| 245467 | 0.67 | 28079099 |
| 245473 | 1.12 | 28079006 |
| 245491 | 0.21 | 28079024 |
| 245495 | 0.67 | 28079099 |

19397 rows × 2 columns

Line chart

```
In [115]: data.plot.line(subplots=True)
```


Out[115]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



Line chart

```
In [116]: data.plot.line()
```

```
Out[116]: <AxesSubplot:>
```




Bar chart

```
In [117]: b=data[0:50]
```

```
In [118]: b.plot.bar()
```


```
Out[118]: <AxesSubplot:>
```



Histogram

```
In [119]: data.plot.hist()
```


```
Out[119]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

```
In [120]: data.plot.area()
```


```
Out[120]: <AxesSubplot:>
```



Box chart

```
In [121]: data.plot.box()
```


```
Out[121]: <AxesSubplot:>
```



Pie chart

```
In [122]: b.plot.pie(y='station' )
```


```
Out[122]: <AxesSubplot:ylabel='station'>
```



Scatter chart

```
In [123]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[123]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [124]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        19397 non-null  object
1   BEN         19397 non-null  float64
2   CO          19397 non-null  float64
3   EBE         19397 non-null  float64
4   MXY         19397 non-null  float64
5   NMHC        19397 non-null  float64
6   NO_2        19397 non-null  float64
7   NOx         19397 non-null  float64
8   OXY         19397 non-null  float64
9   O_3         19397 non-null  float64
10  PM10        19397 non-null  float64
11  PM25        19397 non-null  float64
12  PXY         19397 non-null  float64
13  SO_2        19397 non-null  float64
14  TCH         19397 non-null  float64
15  TOL         19397 non-null  float64
16  station     19397 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 2.7+ MB
```


In [125]:

df.describe()

Out[125]:

| | BEN | CO | EBE | MXY | NMHC | NO_2 | |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 19397.000000 | 19397.000000 | 19397.000000 | 19397.000000 | 19397.000000 | 19397.000000 | 19397.000000 |
| mean | 2.250781 | 0.675347 | 2.775913 | 5.424809 | 0.151024 | 62.887023 | 19397.000000 |
| std | 2.184724 | 0.591026 | 2.729622 | 5.554358 | 0.158603 | 37.952255 | 19397.000000 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.090000 | 19397.000000 |
| 25% | 0.870000 | 0.320000 | 1.020000 | 1.780000 | 0.060000 | 35.150002 | 19397.000000 |
| 50% | 1.620000 | 0.520000 | 1.970000 | 3.800000 | 0.110000 | 58.310001 | 19397.000000 |
| 75% | 2.910000 | 0.860000 | 3.580000 | 7.260000 | 0.200000 | 85.730003 | 19397.000000 |
| max | 34.180000 | 8.900000 | 41.880001 | 91.599998 | 4.810000 | 355.100006 | 19397.000000 |

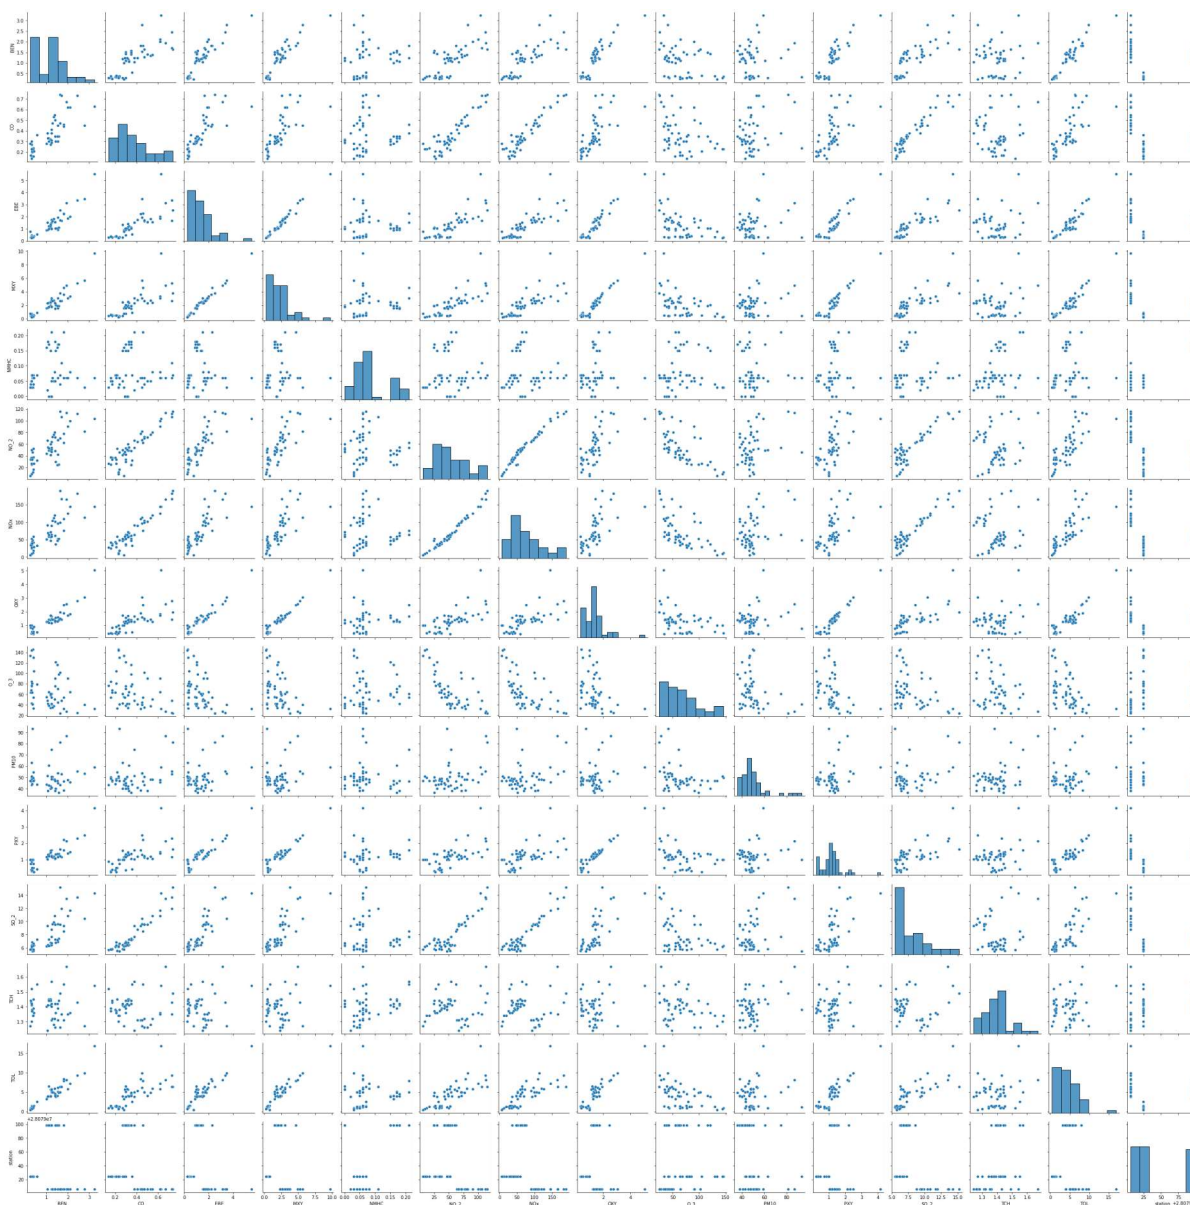
In [126]:

df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]

EDA AND VISUALIZATION

```
In [138]: sns.pairplot(df1[0:50])
```

```
Out[138]: <seaborn.axisgrid.PairGrid at 0x1cd2081d5e0>
```




```
In [139]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[139]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [140]: sns.heatmap(df1.corr())
```

```
Out[140]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [141]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [142]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [143]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[143]: LinearRegression()

```
In [144]: lr.intercept_
```

Out[144]: 28079074.904809926


```
In [145]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[145]:

| | Co-efficient |
|-------------|--------------|
| BEN | -4.149320 |
| CO | 28.974299 |
| EBE | 4.093538 |
| MXY | -3.564947 |
| NMHC | 76.214166 |
| NO_2 | -0.152042 |
| NOx | -0.263485 |
| OXY | -2.700986 |
| O_3 | -0.299274 |
| PM10 | 0.091041 |
| PXY | 6.509818 |
| SO_2 | -0.207503 |
| TCH | -6.417410 |
| TOL | 1.237076 |

```
In [146]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[146]: <matplotlib.collections.PathCollection at 0x1cd000e0c40>



ACCURACY

```
In [147]: lr.score(x_test,y_test)
```

Out[147]: 0.10368510261153874

```
In [148]: lr.score(x_train,y_train)
```

Out[148]: 0.10720468459527421

Ridge and Lasso

```
In [149]: from sklearn.linear_model import Ridge,Lasso
```

```
In [150]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[150]: Ridge(alpha=10)

Accuracy(Ridge)

```
In [151]: rr.score(x_test,y_test)
```

```
Out[151]: 0.1013285735884294
```

```
In [152]: rr.score(x_train,y_train)
```

```
Out[152]: 0.10691065764218544
```

```
In [153]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[153]: Lasso(alpha=10)
```

```
In [154]: la.score(x_train,y_train)
```

```
Out[154]: 0.0562981891575286
```

Accuracy(Lasso)

```
In [155]: la.score(x_test,y_test)
```

```
Out[155]: 0.04747300651618447
```

```
In [156]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[156]: ElasticNet()
```

```
In [157]: en.coef_
```

```
Out[157]: array([-0.          ,  0.42069669,  1.41170788, -1.88395165,  0.          ,  
                -0.16846589, -0.09328818, -0.          , -0.23225802,  0.11512537,  
                0.38645125, -0.13953972,  0.          ,  1.21346959])
```

```
In [158]: en.intercept_
```

```
Out[158]: 28079067.29497689
```

```
In [159]: prediction=en.predict(x_test)
```

```
In [160]: en.score(x_test,y_test)
```

```
Out[160]: 0.06055311810140007
```

Evaluation Metrics

```
In [161]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
38.63234705078924
1666.5565468600973
40.823480337424655
```

Logistic Regression

```
In [162]: from sklearn.linear_model import LogisticRegression
```

```
In [163]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [164]: feature_matrix.shape
```

```
Out[164]: (19397, 14)
```

```
In [165]: target_vector.shape
```

```
Out[165]: (19397,)
```

```
In [166]: from sklearn.preprocessing import StandardScaler
```

```
In [167]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [168]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[168]: LogisticRegression(max_iter=10000)
```

```
In [169]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [170]: prediction=logr.predict(observation)
          print(prediction)
```

```
[28079006]
```

```
In [171]: logr.classes_
```

```
Out[171]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [172]: logr.score(fs,target_vector)
```

```
Out[172]: 0.7360416559261741
```

```
In [173]: logr.predict_proba(observation)[0][0]
```

```
Out[173]: 0.9999978255573396
```

```
In [174]: logr.predict_proba(observation)
```

```
Out[174]: array([[9.99997826e-01, 7.75018107e-20, 2.17444266e-06]])
```

Random Forest

```
In [175]: from sklearn.ensemble import RandomForestClassifier
```

```
In [176]: rfc=RandomForestClassifier()
          rfc.fit(x_train,y_train)
```

```
Out[176]: RandomForestClassifier()
```

```
In [177]: parameters={'max_depth':[1,2,3,4,5],
                      'min_samples_leaf':[5,10,15,20,25],
                      'n_estimators':[10,20,30,40,50]}
          }
```



```
In [178]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
grid_search.fit(x_train, y_train)
```

```
Out[178]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [179]: grid_search.best_score_
```

```
Out[179]: 0.7760189647834945
```

```
In [180]: rfc_best = grid_search.best_estimator_
```

```
In [181]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b', 'c'],
```

```
Text(2539.8620689655177, 181.19999999999982, 'gini = 0.349\nsamples = 1250\nvalue = [207, 209, 1601]\nclass = c'),
Text(2924.689655172414, 906.0, 'EBE <= 4.325\ngini = 0.545\nsamples = 2431\nvalue = [2003, 215, 1552]\nclass = a'),
Text(2770.7586206896553, 543.5999999999999, 'NO_2 <= 89.8\ngini = 0.541\nsamples = 2293\nvalue = [1953, 207, 1389]\nclass = a'),
Text(2693.7931034482763, 181.19999999999982, 'gini = 0.551\nsamples = 1700\nvalue = [1298, 149, 1184]\nclass = a'),
Text(2847.724137931035, 181.19999999999982, 'gini = 0.437\nsamples = 593\nvalue = [655, 58, 205]\nclass = a'),
Text(3078.6206896551726, 543.5999999999999, 'O_3 <= 6.98\ngini = 0.404\nsamples = 138\nvalue = [50, 8, 163]\nclass = c'),
Text(3001.6551724137935, 181.19999999999982, 'gini = 0.127\nsamples = 55\nvalue = [6, 0, 82]\nclass = c'),
Text(3155.586206896552, 181.19999999999982, 'gini = 0.516\nsamples = 83\nvalue = [44, 8, 81]\nclass = c'),
Text(3848.275862068966, 1268.4, 'NOx <= 143.5\ngini = 0.39\nsamples = 1752\nvalue = [2091, 25, 714]\nclass = a'),
Text(3540.4137931034484, 906.0, 'TCH <= 1.365\ngini = 0.497\nsamples = 170\nvalue = [113, 5, 171]\nclass = c'),
```

Conclusion

Accuracy

Linear Regression: 0.11045993310581825

Ridge Regression:0.11011845641033136

Lasso Regression:0.044541624988104433

ElasticNet Regression:0.055622313951294466

Logistic Regression:0.7360416559261741

Random Forest:0.7754293098484298

From the above data, we can conclude that logistic regression and random forest is preferable to other regression types

In []: