

## Importing Libraries

```
In [182]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

## Importing Datasets

```
In [305]: df=pd.read_csv(r"C:\Users\user\Desktop\csvs_per_year\csvs_per_year\madrid_2007.
df
```

Out[305]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2007-12-01 01:00:00	NaN	2.86	NaN	NaN	NaN	282.200012	1054.000000	NaN	4.030000	156.1
1	2007-12-01 01:00:00	NaN	1.82	NaN	NaN	NaN	86.419998	354.600006	NaN	3.260000	80.8
2	2007-12-01 01:00:00	NaN	1.47	NaN	NaN	NaN	94.639999	319.000000	NaN	5.310000	53.0
3	2007-12-01 01:00:00	NaN	1.64	NaN	NaN	NaN	127.900002	476.700012	NaN	4.500000	105.3
4	2007-12-01 01:00:00	4.64	1.86	4.26	7.98	0.57	145.100006	573.900024	3.49	52.689999	106.5
...	...	...	...	...	...	...	...	...	...	...	...
225115	2007-03-01 00:00:00	0.30	0.45	1.00	0.30	0.26	8.690000	11.690000	1.00	42.209999	6.7
225116	2007-03-01 00:00:00	NaN	0.16	NaN	NaN	NaN	46.820000	51.480000	NaN	22.150000	5.7
225117	2007-03-01 00:00:00	0.24	NaN	0.20	NaN	0.09	51.259998	66.809998	NaN	18.540001	13.0
225118	2007-03-01 00:00:00	0.11	NaN	1.00	NaN	0.05	24.240000	36.930000	NaN	NaN	6.6
225119	2007-03-01 00:00:00	0.53	0.40	1.00	1.70	0.12	32.360001	47.860001	1.37	24.150000	10.2

225120 rows × 17 columns

# Data Cleaning and Data Preprocessing

```
In [306]: df=df.dropna()
```

```
In [307]: df.columns
```

```
Out[307]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
                'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
              dtype='object')
```

```
In [308]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 25443 entries, 4 to 225119  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   date        25443 non-null  object  
1   BEN         25443 non-null  float64  
2   CO          25443 non-null  float64  
3   EBE         25443 non-null  float64  
4   MXY         25443 non-null  float64  
5   NMHC        25443 non-null  float64  
6   NO_2        25443 non-null  float64  
7   NOx         25443 non-null  float64  
8   OXY         25443 non-null  float64  
9   O_3         25443 non-null  float64  
10  PM10        25443 non-null  float64  
11  PM25        25443 non-null  float64  
12  PXY         25443 non-null  float64  
13  SO_2        25443 non-null  float64  
14  TCH         25443 non-null  float64  
15  TOL         25443 non-null  float64  
16  station     25443 non-null  int64  
dtypes: float64(15), int64(1), object(1)  
memory usage: 3.5+ MB
```

```
In [309]: data=df[['CO' , 'station']]
data
```

Out[309]:


	CO	station
4	1.86	28079006
21	0.31	28079024
25	1.42	28079099
30	1.89	28079006
47	0.30	28079024
...	...	...
225073	0.47	28079006
225094	0.45	28079099
225098	0.41	28079006
225115	0.45	28079024
225119	0.40	28079099

25443 rows × 2 columns

## Line chart

```
In [310]: data.plot.line(subplots=True)
```


Out[310]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



## Line chart

```
In [311]: data.plot.line()
```

```
Out[311]: <AxesSubplot:>
```




## Bar chart

```
In [312]: b=data[0:50]
```

```
In [313]: b.plot.bar()
```


```
Out[313]: <AxesSubplot:>
```



## Histogram

```
In [314]: data.plot.hist()
```


```
Out[314]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [315]: data.plot.area()
```


```
Out[315]: <AxesSubplot:>
```



## Box chart

```
In [316]: data.plot.box()
```


```
Out[316]: <AxesSubplot:>
```



## Pie chart

```
In [317]: b.plot.pie(y='station' )
```


Out[317]: <AxesSubplot:ylabel='station'>



# Scatter chart

```
In [318]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[318]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [319]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25443 entries, 4 to 225119
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   date        25443 non-null  object  
 1   BEN         25443 non-null  float64 
 2   CO          25443 non-null  float64 
 3   EBE         25443 non-null  float64 
 4   MXY         25443 non-null  float64 
 5   NMHC        25443 non-null  float64 
 6   NO_2        25443 non-null  float64 
 7   NOx         25443 non-null  float64 
 8   OXY         25443 non-null  float64 
 9   O_3         25443 non-null  float64 
10  PM10        25443 non-null  float64 
11  PM25        25443 non-null  float64 
12  PXY         25443 non-null  float64 
13  SO_2        25443 non-null  float64 
14  TCH         25443 non-null  float64 
15  TOL         25443 non-null  float64 
16  station     25443 non-null  int64   
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```



In [320]: `df.describe()`

Out[320]:

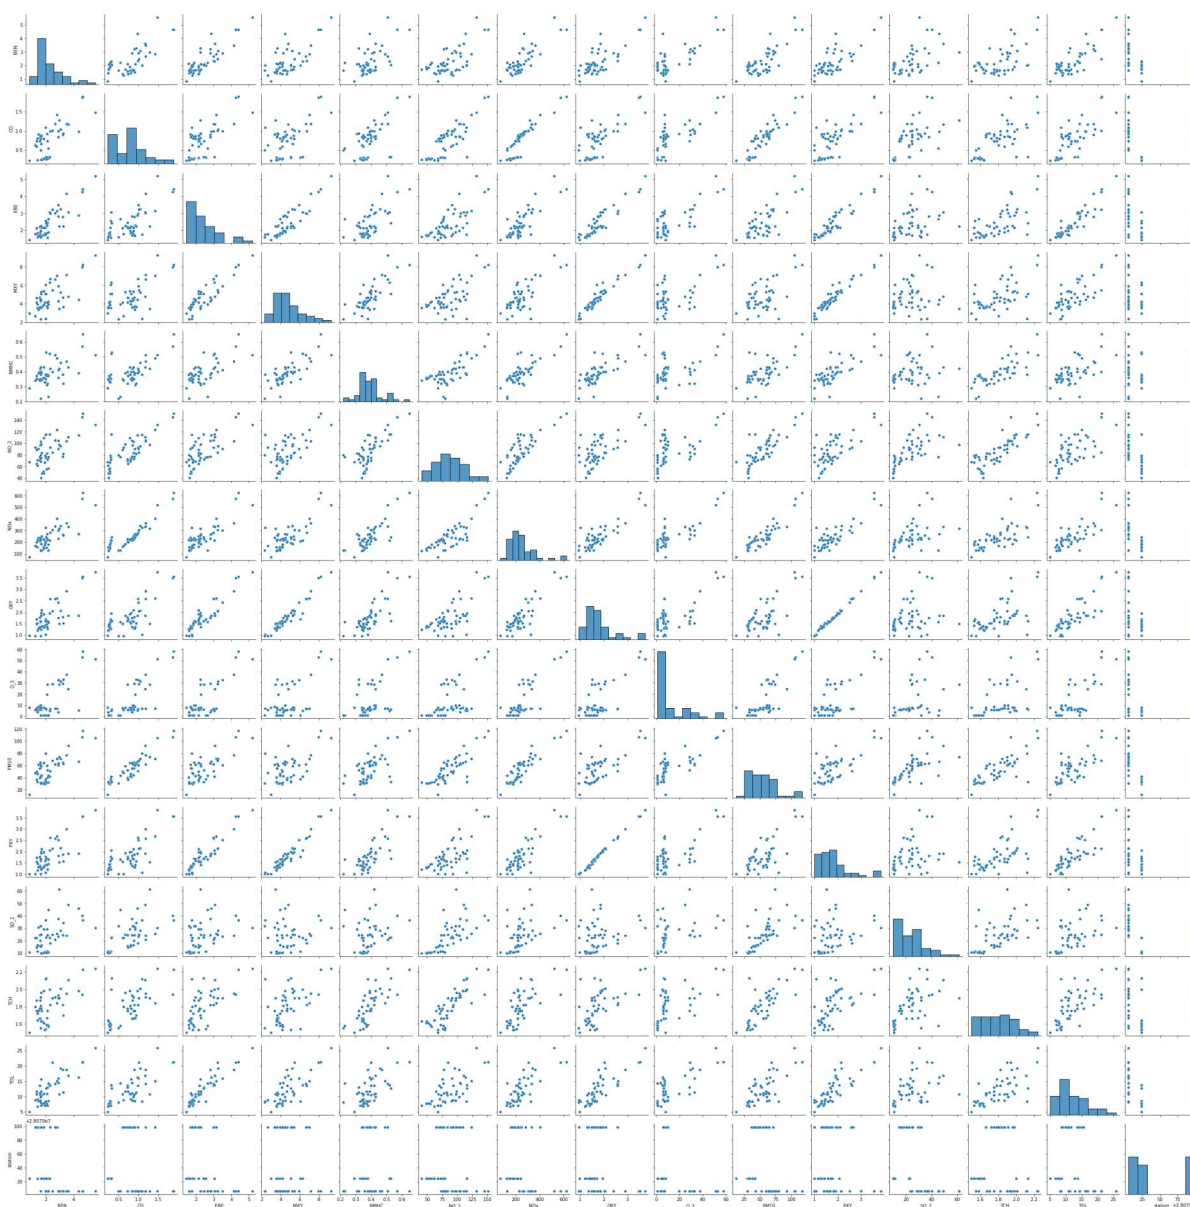
	BEN	CO	EBE	MXY	NMHC	NO_2	
<b>count</b>	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000
<b>mean</b>	1.146744	0.505120	1.394071	2.392008	0.249967	58.532683	1.146744
<b>std</b>	1.278733	0.423231	1.268265	2.784302	0.142627	37.755029	1.278733
<b>min</b>	0.130000	0.000000	0.120000	0.150000	0.000000	1.690000	0.130000
<b>25%</b>	0.450000	0.260000	0.780000	0.960000	0.160000	31.285001	0.450000
<b>50%</b>	0.770000	0.400000	1.000000	1.500000	0.220000	54.080002	0.770000
<b>75%</b>	1.390000	0.640000	1.580000	2.855000	0.300000	79.230003	1.390000
<b>max</b>	30.139999	9.660000	31.680000	65.480003	2.570000	430.299988	30.139999

In [321]: `df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`

## EDA AND VISUALIZATION

```
In [322]: sns.pairplot(df1[0:50])
```

```
Out[322]: <seaborn.axisgrid.PairGrid at 0x1cd3f6a0370>
```




```
In [323]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[323]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [324]: sns.heatmap(df1.corr())
```

```
Out[324]: <AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [325]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [326]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [327]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[327]: LinearRegression()

```
In [328]: lr.intercept_
```

Out[328]: 28079008.864636354


```
In [329]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[329]:

	Co-efficient
<b>BEN</b>	-32.236540
<b>CO</b>	18.829784
<b>EBE</b>	0.554598
<b>MXY</b>	-1.207037
<b>NMHC</b>	-41.467572
<b>NO_2</b>	0.132824
<b>NOx</b>	-0.054933
<b>OXY</b>	3.444702
<b>O_3</b>	-0.030696
<b>PM10</b>	0.148554
<b>PXY</b>	9.242893
<b>SO_2</b>	0.194962
<b>TCH</b>	26.206990
<b>TOL</b>	3.044746

```
In [330]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[330]: <matplotlib.collections.PathCollection at 0x1cd630c5490>



## ACCURACY

```
In [331]: lr.score(x_test,y_test)
```

Out[331]: 0.1656836298520762

```
In [332]: lr.score(x_train,y_train)
```

Out[332]: 0.15576734872658937

## Ridge and Lasso

```
In [333]: from sklearn.linear_model import Ridge,Lasso
```

```
In [334]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[334]: Ridge(alpha=10)

## Accuracy(Ridge)

```
In [335]: rr.score(x_test,y_test)
```

```
Out[335]: 0.16576377535882858
```

```
In [336]: rr.score(x_train,y_train)
```

```
Out[336]: 0.15571427620227551
```

```
In [337]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[337]: Lasso(alpha=10)
```

```
In [338]: la.score(x_train,y_train)
```

```
Out[338]: 0.013396609870542253
```

## Accuracy(Lasso)

```
In [339]: la.score(x_test,y_test)
```

```
Out[339]: 0.013655657610385008
```

```
In [340]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[340]: ElasticNet()
```

```
In [341]: en.coef_
```

```
Out[341]: array([-7.88215711,  0.          ,  0.          ,  0.          , -0.          ,  
                0.06218596, -0.05698424,  0.51989087, -0.05151554,  0.16619749,  
                0.58618021,  0.          ,  0.          ,  1.058272  ])
```

```
In [342]: en.intercept_
```

```
Out[342]: 28079045.65718712
```

```
In [343]: prediction=en.predict(x_test)
```

```
In [344]: en.score(x_test,y_test)
```

```
Out[344]: 0.06828925225659888
```

## Evaluation Metrics

```
In [345]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
36.59806709897106
1531.6968382280866
39.13689867922708
```

## Logistic Regression

```
In [346]: from sklearn.linear_model import LogisticRegression
```

```
In [347]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [348]: feature_matrix.shape
```

```
Out[348]: (25443, 14)
```

```
In [349]: target_vector.shape
```

```
Out[349]: (25443,)
```

```
In [350]: from sklearn.preprocessing import StandardScaler
```

```
In [351]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [352]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[352]: LogisticRegression(max_iter=10000)
```

```
In [353]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [354]: prediction=logr.predict(observation)
          print(prediction)
```

```
[28079099]
```

```
In [355]: logr.classes_
```

```
Out[355]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [356]: logr.score(fs,target_vector)
```

```
Out[356]: 0.8146838030106512
```

```
In [357]: logr.predict_proba(observation)[0][0]
```

```
Out[357]: 1.082753977181323e-19
```

```
In [358]: logr.predict_proba(observation)
```

```
Out[358]: array([[1.08275398e-19, 1.80383815e-19, 1.00000000e+00]])
```

## Random Forest

```
In [359]: from sklearn.ensemble import RandomForestClassifier
```

```
In [360]: rfc=RandomForestClassifier()
          rfc.fit(x_train,y_train)
```

```
Out[360]: RandomForestClassifier()
```

```
In [361]: parameters={'max_depth':[1,2,3,4,5],
                      'min_samples_leaf':[5,10,15,20,25],
                      'n_estimators':[10,20,30,40,50]}
          }
```



```
In [362]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
grid_search.fit(x_train, y_train)
```

```
Out[362]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [363]: grid_search.best_score_
```

```
Out[363]: 0.8251544076361594
```

```
In [364]: rfc_best = grid_search.best_estimator_
```

```
In [365]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b', 'c'],
          value=[0, 41, 0]\n\nclass = b'),
          Text(1370.5263157894735, 1268.4, 'OXY <= 0.365\ngini = 0.573\nsamples = 2389\nvalue = [606, 2140, 960]\n\nclass = b'),
          Text(1096.421052631579, 906.0, 'NOx <= 18.98\ngini = 0.172\nsamples = 127\nvalue = [193, 18, 2]\n\nclass = a'),
          Text(1018.1052631578947, 543.5999999999999, 'gini = 0.5\nsamples = 7\nvalue = [5, 5, 0]\n\nclass = a'),
          Text(1174.7368421052631, 543.5999999999999, 'TCH <= 1.385\ngini = 0.138\nsamples = 120\nvalue = [188, 13, 2]\n\nclass = a'),
          Text(1096.421052631579, 181.19999999999982, 'gini = 0.063\nsamples = 109\nvalue = [179, 4, 2]\n\nclass = a'),
          Text(1253.0526315789473, 181.19999999999982, 'gini = 0.5\nsamples = 11\nvalue = [9, 9, 0]\n\nclass = a'),
          Text(1644.6315789473683, 906.0, 'NMHC <= 0.225\ngini = 0.542\nsamples = 2262\nvalue = [413, 2122, 958]\n\nclass = b'),
          Text(1488.0, 543.5999999999999, 'SO_2 <= 5.225\ngini = 0.613\nsamples = 1071\nvalue = [331, 487, 869]\n\nclass = c'),
          Text(1409.6842105263156, 181.19999999999982, 'gini = 0.381\nsamples = 270\nvalue = [63, 321, 34]\n\nclass = b'),
          Text(1566.3157894736842, 181.19999999999982, 'gini = 0.505\nsamples = 801\nvalue = [566, 466, 269]\n\nclass = a'),
```

## Conclusion

## Accuracy

Linear Regression: 0.16331457098631952

Ridge Regression:0.16317654437433604

Lasso Regression:0.013732764982463452

ElasticNet Regression:0.0693172677037851

Logistic Regression:0.8146838030106512

Random Forest:0.8748413156376227

**From the above data, we can conclude that random forest is preferable to other regression types**

In [ ]: