

Importing Libraries

```
In [428]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Datasets

```
In [490]: df=pd.read_csv(r"C:\Users\user\Desktop\csvs_per_year\csvs_per_year\madrid_2009.
df
```

Out[490]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PM
0	2009-10-01 01:00:00	NaN	0.27	NaN	NaN	NaN	39.889999	48.150002	NaN	50.680000	18.2600
1	2009-10-01 01:00:00	NaN	0.22	NaN	NaN	NaN	21.230000	24.260000	NaN	55.880001	10.5800
2	2009-10-01 01:00:00	NaN	0.18	NaN	NaN	NaN	31.230000	34.880001	NaN	49.060001	25.1900
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.5300
4	2009-10-01 01:00:00	NaN	0.41	NaN	NaN	0.12	61.349998	76.260002	NaN	38.090000	23.7600
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.8300
215684	2009-06-01 00:00:00	NaN	0.31	NaN	NaN	NaN	76.110001	101.099998	NaN	41.220001	9.9200
215685	2009-06-01 00:00:00	0.13	NaN	0.86	NaN	0.23	81.050003	99.849998	NaN	24.830000	12.4600
215686	2009-06-01 00:00:00	0.21	NaN	2.96	NaN	0.10	72.419998	82.959999	NaN	NaN	13.0300
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.3600

215688 rows × 17 columns

Data Cleaning and Data Preprocessing

```
In [430]: df=df.dropna()
```

```
In [431]: df.columns
```

```
Out[431]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
                'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
              dtype='object')
```

```
In [432]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 24717 entries, 3 to 215687  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   date        24717 non-null  object  
1   BEN         24717 non-null  float64  
2   CO          24717 non-null  float64  
3   EBE         24717 non-null  float64  
4   MXY         24717 non-null  float64  
5   NMHC        24717 non-null  float64  
6   NO_2        24717 non-null  float64  
7   NOx         24717 non-null  float64  
8   OXY         24717 non-null  float64  
9   O_3         24717 non-null  float64  
10  PM10        24717 non-null  float64  
11  PM25        24717 non-null  float64  
12  PXY         24717 non-null  float64  
13  SO_2        24717 non-null  float64  
14  TCH         24717 non-null  float64  
15  TOL         24717 non-null  float64  
16  station     24717 non-null  int64  
dtypes: float64(15), int64(1), object(1)  
memory usage: 3.4+ MB
```

```
In [433]: data=df[['CO' , 'station']]
data
```

Out[433]:


	CO	station
3	0.33	28079006
20	0.32	28079024
24	0.24	28079099
28	0.21	28079006
45	0.30	28079024
...
215659	0.27	28079024
215663	0.35	28079099
215667	0.29	28079006
215683	0.22	28079024
215687	0.32	28079099

24717 rows × 2 columns

Line chart

```
In [434]: data.plot.line(subplots=True)
```


Out[434]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)



Line chart

```
In [435]: data.plot.line()
```

```
Out[435]: <AxesSubplot:>
```




Bar chart

```
In [436]: b=data[0:50]
```

```
In [437]: b.plot.bar()
```


```
Out[437]: <AxesSubplot:>
```



Histogram

```
In [438]: data.plot.hist()
```


```
Out[438]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

```
In [439]: data.plot.area()
```


```
Out[439]: <AxesSubplot:>
```



Box chart

```
In [440]: data.plot.box()
```


```
Out[440]: <AxesSubplot:>
```



Pie chart

```
In [441]: b.plot.pie(y='station' )
```


```
Out[441]: <AxesSubplot:ylabel='station'>
```



Scatter chart

```
In [442]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[442]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [443]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24717 entries, 3 to 215687
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   date        24717 non-null  object  
1   BEN         24717 non-null  float64 
2   CO          24717 non-null  float64 
3   EBE         24717 non-null  float64 
4   MXY         24717 non-null  float64 
5   NMHC        24717 non-null  float64 
6   NO_2        24717 non-null  float64 
7   NOx         24717 non-null  float64 
8   OXY         24717 non-null  float64 
9   O_3         24717 non-null  float64 
10  PM10        24717 non-null  float64 
11  PM25        24717 non-null  float64 
12  PXY         24717 non-null  float64 
13  SO_2        24717 non-null  float64 
14  TCH         24717 non-null  float64 
15  TOL         24717 non-null  float64 
16  station     24717 non-null  int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```


In [444]: `df.describe()`

Out[444]:

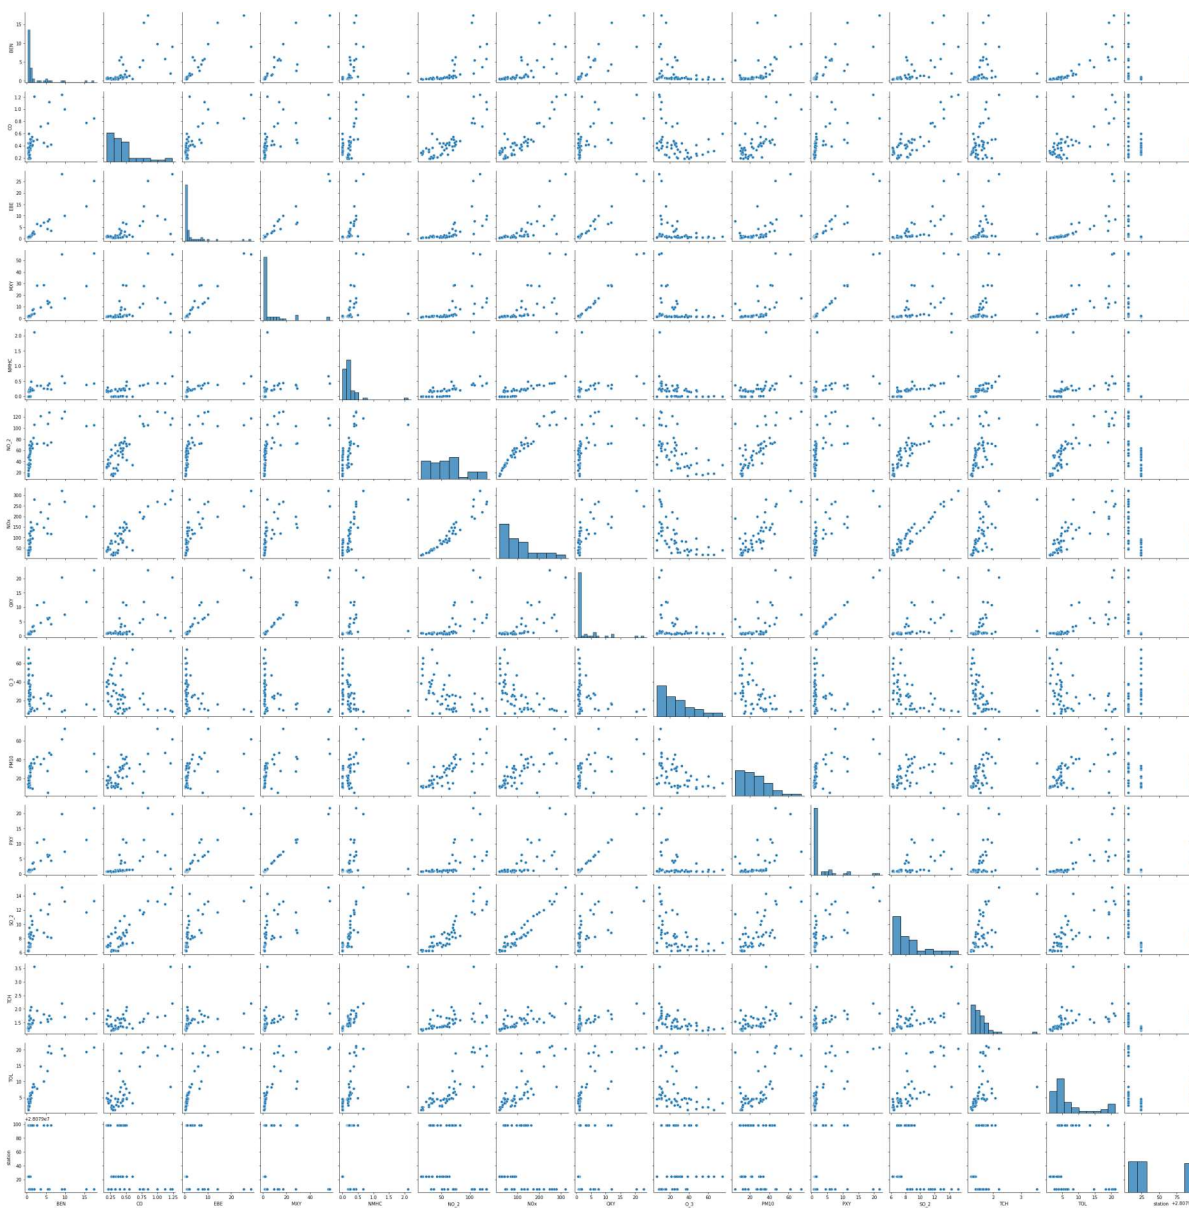
	BEN	CO	EBE	MXY	NMHC	NO_2	
count	24717.000000	24717.000000	24717.000000	24717.000000	24717.000000	24717.000000	247
mean	1.010583	0.448056	1.262430	2.244469	0.219582	55.563929	!
std	1.007345	0.291706	1.074768	2.242214	0.141661	38.911677	!
min	0.170000	0.060000	0.250000	0.240000	0.000000	0.600000	
25%	0.460000	0.270000	0.720000	0.990000	0.140000	26.510000	:
50%	0.670000	0.370000	1.000000	1.490000	0.190000	47.930000	(
75%	1.180000	0.570000	1.430000	2.820000	0.260000	76.269997	1:
max	22.379999	5.570000	47.669998	56.500000	2.580000	477.399994	14:

In [445]: `df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]`

EDA AND VISUALIZATION

```
In [446]: sns.pairplot(df1[0:50])
```

```
Out[446]: <seaborn.axisgrid.PairGrid at 0x1cd915c6730>
```




```
In [447]: sns.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

```
Out[447]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [448]: sns.heatmap(df1.corr())
```

```
Out[448]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [449]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
y=df['station']
```

```
In [450]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [451]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[451]: LinearRegression()

```
In [452]: lr.intercept_
```

Out[452]: 28078891.75325078


```
In [453]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[453]:

	Co-efficient
BEN	-37.039973
CO	-30.754126
EBE	7.531596
MXY	-1.633098
NMHC	-18.185915
NO_2	-0.183246
NOx	0.207947
OXY	14.764541
O_3	0.024889
PM10	-0.051958
PXY	2.828355
SO_2	-0.302665
TCH	124.896713
TOL	-1.148805

```
In [454]: prediction =lr.predict(x_test)
plt.scatter(y_test,prediction)
```

Out[454]: <matplotlib.collections.PathCollection at 0x1cd630e3070>



ACCURACY

```
In [455]: lr.score(x_test,y_test)
```

Out[455]: 0.26828993596392614

```
In [456]: lr.score(x_train,y_train)
```

Out[456]: 0.2946269921837963

Ridge and Lasso

```
In [457]: from sklearn.linear_model import Ridge,Lasso
```

```
In [458]: rr=Ridge(alpha=10)
rr.fit(x_train,y_train)
```

Out[458]: Ridge(alpha=10)

Accuracy(Ridge)

```
In [459]: rr.score(x_test,y_test)
```

```
Out[459]: 0.2700143791476979
```

```
In [460]: rr.score(x_train,y_train)
```

```
Out[460]: 0.2942743131240403
```

```
In [461]: la=Lasso(alpha=10)  
la.fit(x_train,y_train)
```

```
Out[461]: Lasso(alpha=10)
```

```
In [462]: la.score(x_train,y_train)
```

```
Out[462]: 0.033961539134792496
```

Accuracy(Lasso)

```
In [463]: la.score(x_test,y_test)
```

```
Out[463]: 0.03863380766863844
```

```
In [464]: from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[464]: ElasticNet()
```

```
In [465]: en.coef_
```

```
Out[465]: array([-7.07079106, -0.68564137,  0.42393327,  2.12561565, -0.  
                -0.24862351,  0.13491165,  1.24030818, -0.13651085,  0.08671355,  
                1.92554874, -0.71980373,  1.47185506, -1.99024297])
```

```
In [466]: en.intercept_
```

```
Out[466]: 28079063.075263444
```

```
In [467]: prediction=en.predict(x_test)
```

```
In [468]: en.score(x_test,y_test)
```

```
Out[468]: 0.11209585600575955
```

Evaluation Metrics

```
In [469]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
35.84169726560099
1460.280829285163
38.21362099154126
```

Logistic Regression

```
In [470]: from sklearn.linear_model import LogisticRegression
```

```
In [471]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']
```

```
In [472]: feature_matrix.shape
```

```
Out[472]: (24717, 14)
```

```
In [473]: target_vector.shape
```

```
Out[473]: (24717,)
```

```
In [474]: from sklearn.preprocessing import StandardScaler
```

```
In [475]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [476]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[476]: LogisticRegression(max_iter=10000)
```

```
In [477]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [478]: prediction=logr.predict(observation)
          print(prediction)
```

```
[28079099]
```

```
In [479]: logr.classes_
```

```
Out[479]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [480]: logr.score(fs,target_vector)
```

```
Out[480]: 0.8951733624630821
```

```
In [481]: logr.predict_proba(observation)[0][0]
```

```
Out[481]: 5.447205522232353e-13
```

```
In [482]: logr.predict_proba(observation)
```

```
Out[482]: array([[5.44720552e-13, 8.28692830e-44, 1.00000000e+00]])
```

Random Forest

```
In [483]: from sklearn.ensemble import RandomForestClassifier
```

```
In [484]: rfc=RandomForestClassifier()
          rfc.fit(x_train,y_train)
```

```
Out[484]: RandomForestClassifier()
```

```
In [485]: parameters={'max_depth':[1,2,3,4,5],
                      'min_samples_leaf':[5,10,15,20,25],
                      'n_estimators':[10,20,30,40,50]
                      }
```



```
In [486]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
grid_search.fit(x_train, y_train)
```

```
Out[486]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [487]: grid_search.best_score_
```

```
Out[487]: 0.9042250653638224
```

```
In [488]: rfc_best = grid_search.best_estimator_
```

```
In [489]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b'],
          node_ids=[12, 28, 71], node_class='c',
          node_value=[12, 28, 71], node_nclass='c'),
          node_ids=[846, 620, 81, 3], node_class='b'),
          node_value=[846, 620, 81, 3], node_nclass='b'),
          node_ids=[1077, 517, 413, 793, 105], node_class='b'),
          node_value=[1077, 517, 413, 793, 105], node_nclass='b'),
          node_ids=[1000, 551, 724, 1379, 312], node_class='b'),
          node_value=[1000, 551, 724, 1379, 312], node_nclass='b'),
          node_ids=[1154, 482, 758, 620, 6898], node_class='b'),
          node_value=[1154, 482, 758, 620, 6898], node_nclass='b'),
          node_ids=[1731, 724, 1379, 310, 346], node_class='c'),
          node_value=[1731, 724, 1379, 310, 346], node_nclass='c'),
          node_ids=[1539, 310, 344, 827, 5863], node_class='c'),
          node_value=[1539, 310, 344, 827, 5863], node_nclass='c'),
          node_ids=[1385, 379, 310, 344, 8277], node_class='b'),
          node_value=[1385, 379, 310, 344, 8277], node_nclass='b'),
          node_ids=[1308, 413, 793, 103, 4484], node_class='b'),
          node_value=[1308, 413, 793, 103, 4484], node_nclass='b'),
          node_ids=[1462, 344, 827, 586, 207], node_class='b'),
          node_value=[1462, 344, 827, 586, 207], node_nclass='b'),
```

Conclusion

Accuracy

Linear Regression: 0.16331457098631952

Ridge Regression:0.16317654437433604

Lasso Regression:0.013732764982463452

ElasticNet Regression:0.0693172677037851

Logistic Regression:0.8951733624630821

Random Forest:0.8911047204272553

From the above data, we can conclude that logistic regression is preferable to other regression types

In []: